

ML2 Takehome project

Implementation of Lenet5 with Libtorch

Anthony Jung

Contents

- Libtorch Lenet5
 - How to define custom module properly
- Project Overall
 - Torchserve
 - Airflow

Libtorch Lenet5

Lenet5




Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

gradient-based learning applied to document recognition, Yann LeCun et al.

Model

layer	kind	num filter	filter size	stride	padding
c1	conv	6	(5,5)	1	valid
s2	max pooling	-	(2,2)	2	valid
c3	conv	16	(5,5)	1	valid
s4	max pooling	-	(2,2)	2	valid
c5	conv	120	(5,5)	1	valid
f6	dense	84	-	-	-
output	softmax	10	-	-	-

Implementation

```
struct Lenet5Impl : nn::Module
{
public:
    Lenet5Impl()
        : c1(nn::Conv2dOptions(1, 6, /*kernel_size=*/{5, 5}).padding({2, 2})),
          c3(nn::Conv2dOptions(6, 16, /*kernel_size=*/{5, 5})),
          c5(nn::Conv2dOptions(16, 120, /*kernel_size=*/{5, 5})),
          f6(120, 84),
          output(84, 10)
    {
        register_module("c1", c1);
        register_module("c3", c3);
        register_module("c5", c5);
        register_module("f6", f6);
        register_module("output", output);
    }
}
```

Implementation

```
Tensor forward(Tensor x)
{
    x = max_pool2d(relu(c1(x)), {2, 2}, {2, 2});
    x = max_pool2d(relu(c3(x)), {2, 2}, {2, 2});
    x = relu(c5(x));
    x = x.view({x.size(0), -1});
    x = relu(f6(x));
    x = log_softmax(output(x), /*dim=*/1);

    return x;
}

nn::Conv2d c1, c3, c5;
nn::Linear f6, output;
};

TORCH_MODULE(Lenet5);
```

Class Definition

```
Tensor forward(Tensor x)
{
    x = max_pool2d(relu(c1(x)), {2, 2}, {2, 2});
    x = max_pool2d(relu(c3(x)), {2, 2}, {2, 2});
    x = relu(c5(x));
    x = x.view({x.size(0), -1});
    x = relu(f6(x));
    x = log_softmax(output(x), /*dim=*/1);

    return x;
}

nn::Conv2d c1, c3, c5;
nn::Linear f6, output;
};

TORCH_MODULE(Lenet5);
```


MACRO : TORCH_MODULE (pimpl.h)

```
#define TORCH_MODULE_IMPL(Name, ImplType) \
    class Name : public torch::nn::ModuleHolder<ImplType> { /* NOLINT */ \
    public: \
        using torch::nn::ModuleHolder<ImplType>::ModuleHolder; \
        using Impl = ImplType; \
    }

#define TORCH_MODULE(Name) TORCH_MODULE_IMPL(Name, Name##Impl)
```

ModuleHolder Class (pimpl.h)

```
template <typename Contained>
class ModuleHolder : torch::detail::ModuleHolderIndicator {
protected:
    std::shared_ptr<Contained> impl;
public:
    using ContainedType = Contained;

    ModuleHolder() : impl_(default_construct()) {
        static_assert(
            std::is_default_constructible<Contained>::value,
            "You are trying to default construct a module which has "
            "no default constructor. Use = nullptr to give it the empty state "
            "(e.g. `Linear linear = nullptr;` instead of `Linear linear;`)." );
    }
    /* implicit */ ModuleHolder(std::nullptr_t) : impl_(nullptr) {}
```

`std::shared_ptr`(in contrast to `std::unique_ptr`)

`std::shared_ptr` is a smart pointer that retains shared ownership of an object through a pointer. Several `shared_ptr` objects may own the same object.

The object is destroyed and its memory deallocated when either of the following happens:

- the last remaining `shared_ptr` owning the object is destroyed;
- the last remaining `shared_ptr` owning the object is assigned another pointer via `operator=` or `reset()`.

Constructor Initializer

```
struct Lenet5Impl : nn::Module
{
public:
    Lenet5Impl()
        : c1(nn::Conv2dOptions(1, 6, /*kernel_size=*/{5, 5}).padding({2, 2})),
          c3(nn::Conv2dOptions(6, 16, /*kernel_size=*/{5, 5})),
          c5(nn::Conv2dOptions(16, 120, /*kernel_size=*/{5, 5})),
          f6(120, 84),
          output(84, 10)
    {
        register_module("c1", c1);
        register_module("c3", c3);
        register_module("c5", c5);
        register_module("f6", f6);
        register_module("output", output);
    }
}
```

Constructor (pimpl.h)

```
template <
    typename Head,
    typename... Tail,
    typename = typename std::enable_if<
        !(torch::detail::is_module_holder_of<Head, ContainedType>::value &&
            (sizeof...(Tail) == 0))>::type>
explicit ModuleHolder(Head&& head, Tail&&... tail)
    : impl_(new Contained(
        std::forward<Head>(head),
        std::forward<Tail>(tail)...)) {}
```

Register Modules

```
struct Lenet5Impl : nn::Module
{
public:
    Lenet5Impl()
        : c1(nn::Conv2dOptions(1, 6, /*kernel_size=*/{5, 5}).padding({2, 2})),
          c3(nn::Conv2dOptions(6, 16, /*kernel_size=*/{5, 5})),
          c5(nn::Conv2dOptions(16, 120, /*kernel_size=*/{5, 5})),
          f6(120, 84),
          output(84, 10)
    {
        register_module("c1", c1);
        register_module("c3", c3);
        register_module("c5", c5);
        register_module("f6", f6);
        register_module("output", output);
    }
}
```

register_module (module.h)

```
OrderedDict<std::string, std::shared_ptr<Module>>> children;  
...  
template <typename ModuleType>  
std::shared_ptr<ModuleType> Module::register_module(  
    std::string name,  
    std::shared_ptr<ModuleType> module) {  
    TORCH_CHECK(!name.empty(), "Submodule name must not be empty");  
    TORCH_CHECK(  
        name.find('.') == std::string::npos,  
        "Submodule name must not contain a dot (got '",  
        name,  
        "')");  
    auto& base_module = children_.insert(std::move(name), std::move(module));  
    return std::dynamic_pointer_cast<ModuleType>(base_module);  
}
```

Trainer : Torch Way

```
template <typename DataLoader>
void train(
    size_t epoch,
    Lenet5 &model,
    torch::Device device,
    DataLoader &data_loader,
    torch::optim::Optimizer &optimizer,
    size_t dataset_size)
```


Trainer : Torch Way

```
{
    model->train();
    size_t batch_idx = 0;
    for (auto &batch : data_loader)
    {
        auto data = batch.data.to(device),
        targets = batch.target.to(device);
        optimizer.zero_grad();
        auto output = model->forward(data);
        auto loss = torch::nll_loss(output,
        targets);
        AT_ASSERT(!std::isnan(loss.template
        item<float>()));
        loss.backward();
        optimizer.step();
    }
}
```

```
if (batch_idx++ % kLogInterval == 0)
{
    std::printf(
        "\rTrain Epoch: %ld [%5ld/%5ld]\n",
        Loss: %.4f",
        epoch,
        batch_idx * batch.data.size(0),
        dataset_size,
        loss.template item<float>());
}
}
```

Tester : Torch Way

```
template <typename DataLoader>
void test(
    Lenet5 &model,
    torch::Device device,
    DataLoader &data_loader,
    size_t dataset_size)
```

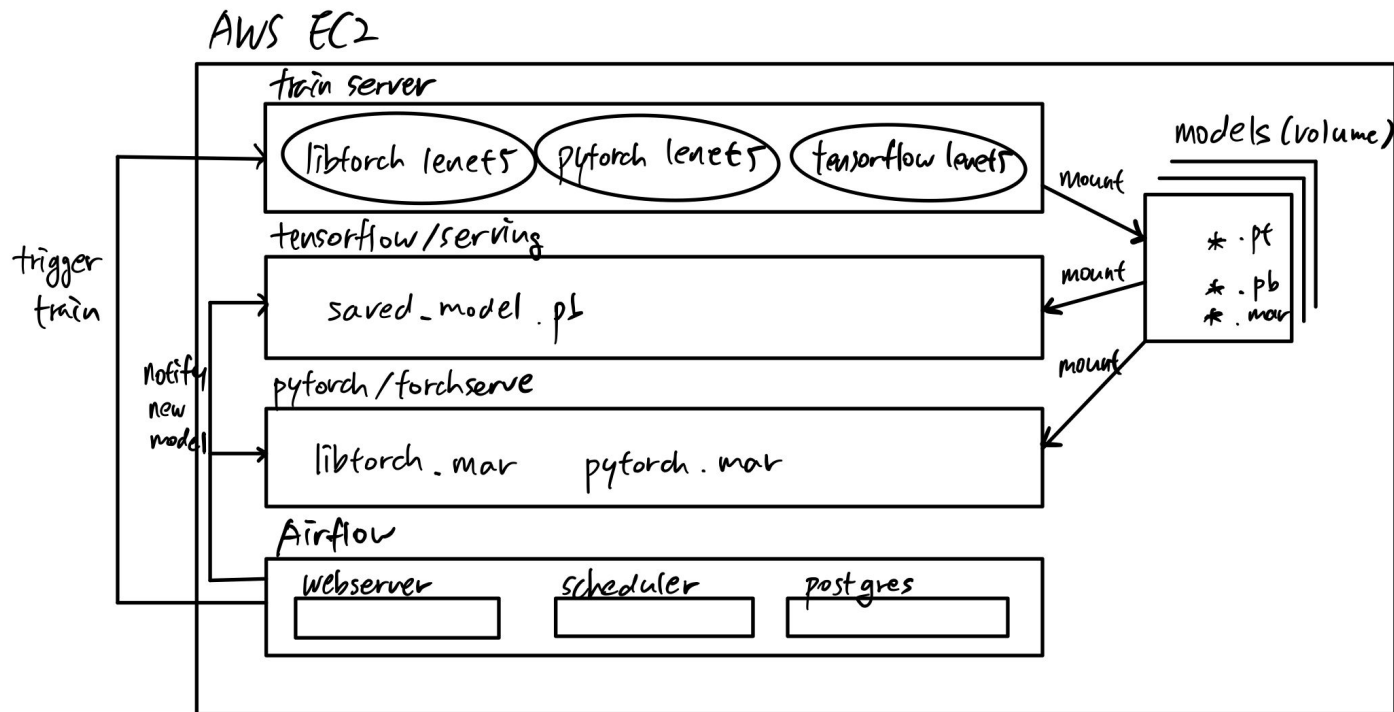
Tester : Torch Way

```
{
    torch::NoGradGuard no_grad;
    model->eval();
    double test_loss = 0;
    int32_t correct = 0;
    for (const auto &batch : data_loader)
    {
        auto data = batch.data.to(device), targets =
batch.target.to(device);
        auto output = model->forward(data);
        test_loss += torch::nll_loss(output, targets,
/*weight=*/{},
torch::Reduction::Sum)
.template item<float>();
        auto pred = output.argmax(1);
        correct += pred.eq(targets).sum().template
item<int64_t>();
    }
```

```
test_loss /= dataset_size;
std::printf(
    "\nTest set: Average loss: %.4f | Accuracy:
%.3f\n",
    test_loss,
    static_cast<double>(correct) / dataset_size);
}
```

Project Overall

Docker Containers



docker-compose

https://github.com/anthony0727/ml2_takehome/blob/master/docker-compose.yml

- run containers
- mount volumes
- CMake
- pull & build docker images
- pip3 install -r requirements.txt
- start server
- ...

API : Tensorflow Serving vs Torchserve

```
tf.keras.models.save_model(  
    model,  
    export_path,  
    overwrite=True,  
    include_optimizer=True,  
    save_format=None,  
    signatures=None,  
    options=None  
)  
signature_def: {  
    key : "my_prediction_signature"  
    value: {  
        inputs: {  
            ...  
        }  
        outputs: {  
            ...  
        }  
        method_name:  
        "tensorflow/serving/predict"  
    }  
}
```

```
class MNISTDigitClassifier(ImageClassifier):  
    """  
    MNISTDigitClassifier handler class. This handler extends  
    class ImageClassifier from image_classifier.py, a  
    default handler. This handler takes an image and returns  
    the number in that image.  
    Here method postprocess() has been overridden while  
    others are reused from parent class.  
    """  
  
    image_processing = transforms.Compose([  
        # transforms.Resize(28),  
        transforms.ToTensor(),  
        transforms.Normalize((0.5,), (0.5,))  
    ])  
  
    def postprocess(self, data):  
        return data.argmax(1).tolist()
```

Facing 503 Internal Error...

```
2020-09-04 12:36:07,574 [INFO ] W-9003-lenet5_1.0-stdout
```

```
org.pytorch.serve.wlm.WorkerLifeCycle -
```

```
torch.nn.modules.module.ModuleAttributeError: 'RecursiveScriptModule' object has  
no attribute 'forward'
```


Torchserve

Model:

Models could be a `script_module` (JIT saved models) or `eager_mode_models`.

These models can provide custom pre- and post-processing of data along with any other model artifacts such as `state_dicts`.

Models can be loaded from cloud storage or from local hosts.

...

Torch::save (serialize.h)

```
template <typename Value, typename... SaveToArgs>
void save(const Value& value, SaveToArgs&&... args) {
    serialize::OutputArchive archive(
        std::make_shared<jit::CompilationUnit>());
    archive << value;
    archive.save_to(std::forward<SaveToArgs>(args)...);
}
```

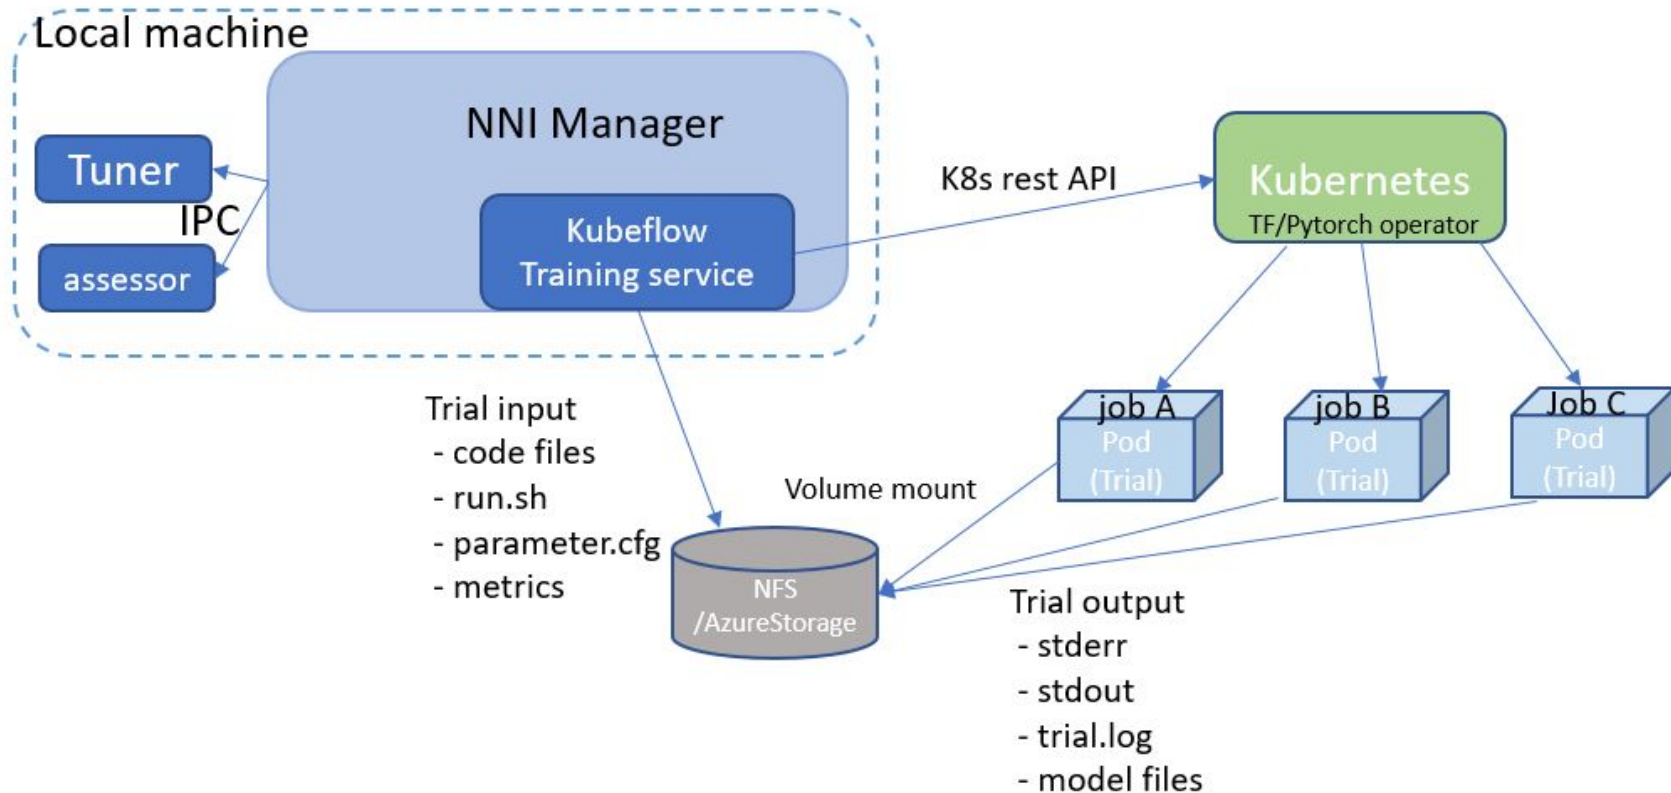
Try tensorflow serving benchmark for now...

https://github.com/anthony0727/ml2_takehome/blob/master/notebook/lenet5_benchmark.ipynb

Application

- Torchserve : web server for many purpose(even for development purpose. '1' server can serve 'n' models)
- Dockerized torchserve container : scalable web service
- Airflow : Continuous transfer learning, failover plan for each task, systemized machine learning platform
- Libtorch : High latency requiring deep learning app

NNI + Kubeflow + Kubernetes(k8s)



Thank you!